


Machine-Level Programming IV: Data

COMP400727: Introduction to Computer Systems

Danfeng Shan
Xi'an Jiaotong University



■ Arrays

- One-dimensional
- Multi-dimensional (nested)
- Multi-level

■ Structures

- Allocation
- Access
- Alignment

■ If we have time: Union

Reminder: Memory Organization

■ Memory locations do not have data types

- Types are implicit in how machine instructions *use* memory

■ Addresses specify byte locations

- Address of a larger datum is the address of its first byte
- Addresses of successive items differ by the item's size

Address	chars	ints	longs
4000			
4001		Addr = 4000	
4002			
4003			Addr = 4000
4004		Addr = 4004	
4005			
4006			
4007			
4008		Addr = 4008	
4009			
400A			Addr = 4008
400B			
400C		Addr = 400C	
400D			
400E			
400F			

Memory Allocation

- C declaration *Type* name [*Length*] ;
 - Array of data type *Type* and length *Length*
 - Contiguously allocated region of $Length * \text{sizeof}(Type)$ bytes in memory

`char string[12];`

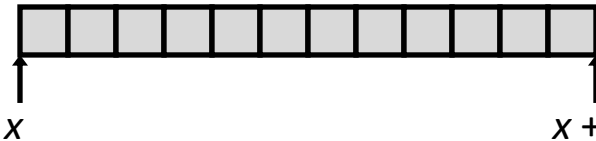


Diagram showing a contiguous memory region for `char string[12]`. The region starts at address x and ends at $x + 12$. The region is divided into 12 small squares, representing individual characters. A red note indicates: "最多可以装11个字符, 因为最后一个是 '\0'" (At most 11 characters can be stored, because the last one is '\0').

`int val[5];`

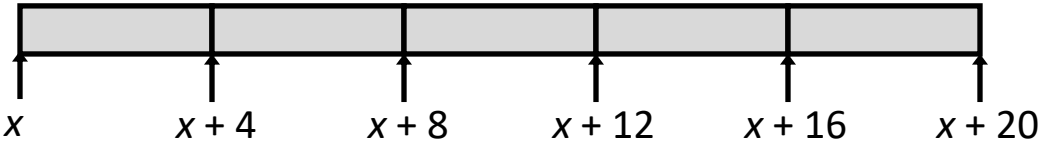


Diagram showing a contiguous memory region for `int val[5]`. The region starts at address x and ends at $x + 20$. The region is divided into 5 larger blocks, representing individual integers. The addresses x , $x + 4$, $x + 8$, $x + 12$, $x + 16$, and $x + 20$ are marked with arrows pointing to the boundaries of the blocks.

`double a[3];`

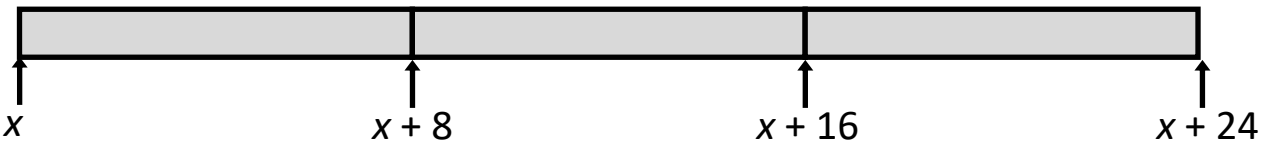


Diagram showing a contiguous memory region for `double a[3]`. The region starts at address x and ends at $x + 24$. The region is divided into 3 very large blocks, representing individual double-precision floating-point numbers. The addresses x , $x + 8$, $x + 16$, and $x + 24$ are marked with arrows pointing to the boundaries of the blocks.

`char *p[3];`
 Typesize = 8

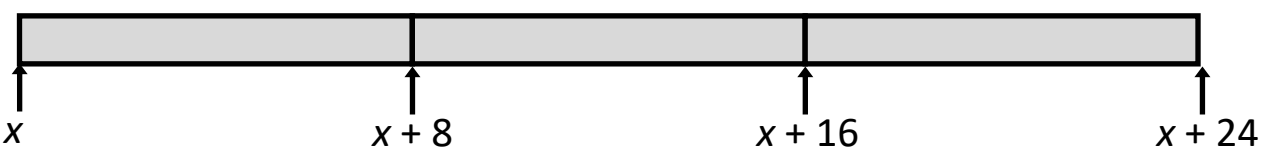



Diagram showing a contiguous memory region for `char *p[3]`. The region starts at address x and ends at $x + 24$. The region is divided into 3 large blocks, representing pointers to character arrays. The addresses x , $x + 8$, $x + 16$, and $x + 24$ are marked with arrows pointing to the boundaries of the blocks. The text "Typesize = 8" is written in red below the declaration.

■ C declaration *Type* *name* [*Length*] ;

- Array of data type *Type* and length *Length*
- Identifier **name** acts like¹ a pointer to array element 0




■ Expression	Type	Value	
<code>val[4]</code>	<code>int</code>	4	
<code>val[6]</code>	<code>int</code>	??	// access past end
<code>val</code>	<code>int *</code>	x	
<code>val+1</code>	<code>int *</code>	$x + 4$	
<code>&val[2]</code>	<code>int *</code>	$x + 8$	// same as <code>val+2</code>
<code>*(val+3)</code>	<code>int</code>	0	// same as <code>val[3]</code>
<code>val + i</code>	<code>int *</code>	$x + 4*i$	// same as <code>&val[i]</code>

¹ in most contexts (but not all)

Example

```
#define ZLEN 6
typedef int zip_dig[ZLEN];


zip_dig xjtu = { 7, 1, 0, 0, 4, 9 };
zip_dig thu  = { 1, 0, 0, 0, 8, 4 };
zip_dig pku  = { 1, 0, 0, 8, 7, 1 };
```



- Declaration “zip_dig xjtu” equivalent to “int xjtu[6]”
- Example arrays were allocated in successive 24 byte blocks
 - Not guaranteed to happen in general

Accessing Example

zip_dig xjtu;



```
int get_digit
(zip_dig z, int digit)
{
    return z[digit];
}
```

x86-64

```
# %rdi = z
# %rsi = digit
movl (%rdi,%rsi,4), %eax # z[digit]
```

基地址 索引号 比例因子

- Register `%rdi` contains starting address of array
- Register `%rsi` contains array index
- Desired digit at $\%rdi + 4 * \%rsi$
- Use memory reference $(\%rdi, \%rsi, 4)$

Array Loop Example

```
void zincr(zip_dig z) {  
    size_t i;  
    for (i = 0; i < ZLEN; i++)  
        z[i]++;  
}
```

```
# %rdi = z  
movl    $0, %eax  
jmp     .L3  
.L4:  
addl    $1, (%rdi,%rax,4)  
addq    $1, %rax  
.L3:  
cmpq    $4, %rax  
jbe     .L4  
rep; ret
```


Array Loop Example

```
void zincr(zip_dig z) {
    size_t i;
    for (i = 0; i < ZLEN; i++)
        z[i]++;
}
```

```
# %rdi = z
movl    $0, %eax           # i = 0
jmp     .L3                # goto middle
.L4:                          # loop:
    addl    $1, (%rdi,%rax,4) # z[i]++
    addq    $1, %rax        # i++
.L3:                          # middle
    cmpq    $5, %rax        # i:4
    jbe     .L4             # if <=, goto loop
rep; ret
```



Understanding Pointers & Arrays #1

Decl	A1 , A2			*A1 , *A2		
	Comp	Bad	Size	Comp	Bad	Size
<code>int A1[3]</code>						
<code>int *A2</code>						

- **Comp: Compiles (Y/N)**
- **Bad: Possible bad pointer reference (Y/N)**
- **Size: Value returned by `sizeof`**

Understanding Pointers & Arrays #1



Decl	A1 , A2			*A1 , *A2		
	Comp	Bad	Size	Comp	Bad	Size
<code>int A1[3]</code>						
<code>int *A2</code>						



- **Comp: Compiles (Y/N)**
- **Bad: Possible bad pointer reference (Y/N)**
- **Size: Value returned by `sizeof`**

Understanding Pointers & Arrays #1



Decl	A1 , A2			*A1 , *A2		
	Comp	Bad	Size	Comp	Bad	Size
<code>int A1[3]</code>	Y	N	12			
<code>int *A2</code>						



- **Comp: Compiles (Y/N)**
- **Bad: Possible bad pointer reference (Y/N)**
- **Size: Value returned by `sizeof`**

Understanding Pointers & Arrays #1



Decl	A1 , A2			*A1 , *A2		
	Comp	Bad	Size	Comp	Bad	Size
<code>int A1[3]</code>	Y	N	12			
<code>int *A2</code>	Y	N	8			



- **Comp: Compiles (Y/N)**
- **Bad: Possible bad pointer reference (Y/N)**
- **Size: Value returned by `sizeof`**

Understanding Pointers & Arrays #1



Decl	A1 , A2			*A1 , *A2		
	Comp	Bad	Size	Comp	Bad	Size
<code>int A1[3]</code>	Y	N	12	Y	N	4
<code>int *A2</code>	Y	N	8			



- **Comp: Compiles (Y/N)**
- **Bad: Possible bad pointer reference (Y/N)**
- **Size: Value returned by `sizeof`**

Understanding Pointers & Arrays #1


Decl	A1 , A2			*A1 , *A2		
	Comp	Bad	Size	Comp	Bad	Size
<code>int A1[3]</code>	Y	N	12	Y	N	4
<code>int *A2</code>	Y	N	8	Y	Y	4



- **Comp: Compiles (Y/N)**
- **Bad: Possible bad pointer reference (Y/N)**
- **Size: Value returned by `sizeof`**


Understanding Pointers & Arrays #2

Decl	<i>An</i>			<i>*An</i>			<i>**An</i>		
	Cmp	Bad	Size	Cmp	Bad	Size	Cmp	Bad	Size
<code>int A1[3]</code>									
<code>int *A2[3]</code>									
<code>int (*A3)[3]</code>									




Understanding Pointers & Arrays #2

Decl	<i>An</i>			<i>*An</i>			<i>**An</i>		
	Cmp	Bad	Size	Cmp	Bad	Size	Cmp	Bad	Size
<code>int A1[3]</code>	Y	N	12						
<code>int *A2[3]</code>									
<code>int (*A3)[3]</code>									




Understanding Pointers & Arrays #2

Decl	<i>An</i>			<i>*An</i>			<i>**An</i>		
	Cmp	Bad	Size	Cmp	Bad	Size	Cmp	Bad	Size
<code>int A1[3]</code>	Y	N	12						
<code>int *A2[3]</code>	Y	N	24						
<code>int (*A3)[3]</code>									




Understanding Pointers & Arrays #2

Decl	<i>An</i>			<i>*An</i>			<i>**An</i>		
	Cmp	Bad	Size	Cmp	Bad	Size	Cmp	Bad	Size
<code>int A1[3]</code>	Y	N	12						
<code>int *A2[3]</code>	Y	N	24						
<code>int (*A3)[3]</code>	Y	N	8						




Understanding Pointers & Arrays #2

Decl	<i>An</i>			<i>*An</i>			<i>**An</i>		
	Cmp	Bad	Size	Cmp	Bad	Size	Cmp	Bad	Size
<code>int A1[3]</code>	Y	N	12	Y	N	4			
<code>int *A2[3]</code>	Y	N	24						
<code>int (*A3)[3]</code>	Y	N	8						



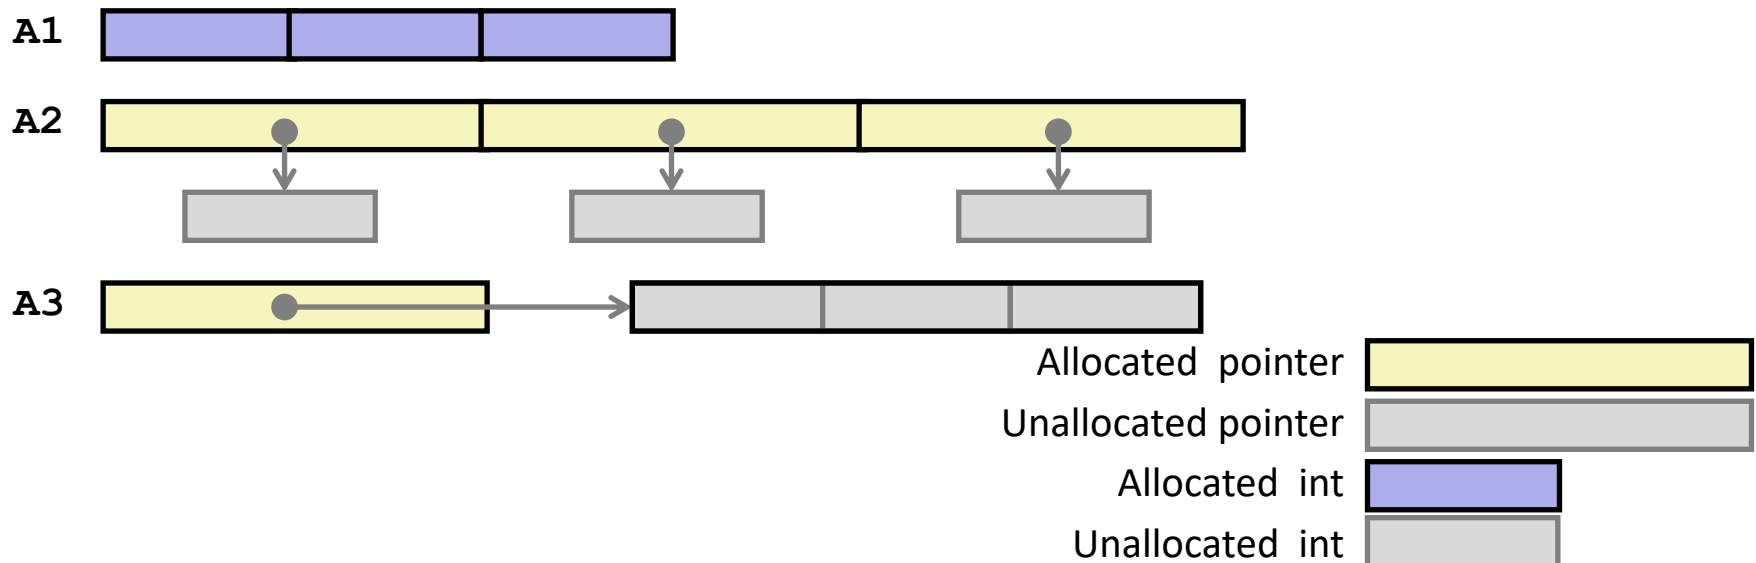
Understanding Pointers & Arrays #2

Decl	<i>An</i>			<i>*An</i>			<i>**An</i>		
	Cmp	Bad	Size	Cmp	Bad	Size	Cmp	Bad	Size
<code>int A1[3]</code>	Y	N	12	Y	N	4			
<code>int *A2[3]</code>	Y	N	24	Y	N	8			
<code>int (*A3)[3]</code>	Y	N	8						



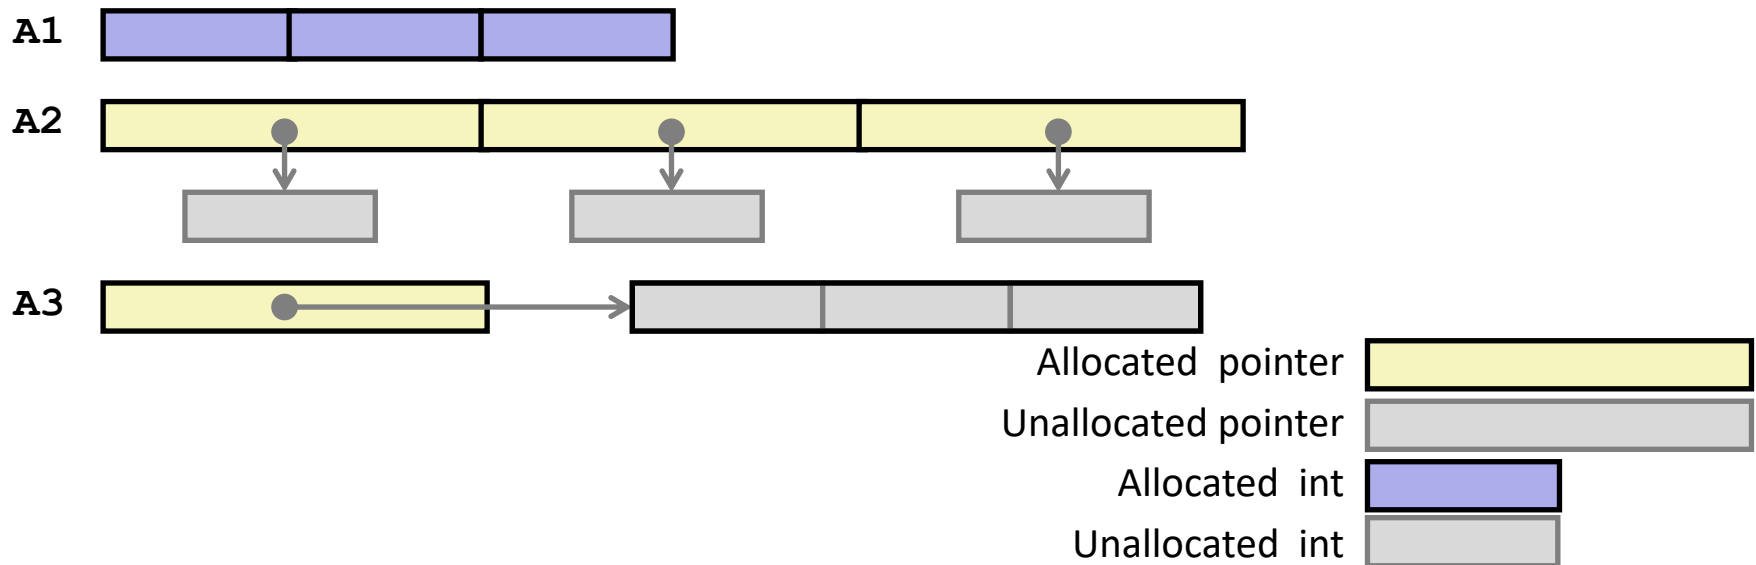
Understanding Pointers & Arrays #2

Decl	<i>An</i>			<i>*An</i>			<i>**An</i>		
	Cmp	Bad	Size	Cmp	Bad	Size	Cmp	Bad	Size
<code>int A1[3]</code>	Y	N	12	Y	N	4			
<code>int *A2[3]</code>	Y	N	24	Y	N	8			
<code>int (*A3)[3]</code>	Y	N	8	Y	Y	12			




Understanding Pointers & Arrays #2

Decl	<i>An</i>			<i>*An</i>			<i>**An</i>		
	Cmp	Bad	Size	Cmp	Bad	Size	Cmp	Bad	Size
<code>int A1[3]</code>	Y	N	12	Y	N	4	N	-	-
<code>int *A2[3]</code>	Y	N	24	Y	N	8			
<code>int (*A3)[3]</code>	Y	N	8	Y	Y	12			




Understanding Pointers & Arrays #2

Decl	<i>An</i>			<i>*An</i>			<i>**An</i>		
	Cmp	Bad	Size	Cmp	Bad	Size	Cmp	Bad	Size
<code>int A1[3]</code>	Y	N	12	Y	N	4	N	-	-
<code>int *A2[3]</code>	Y	N	24	Y	N	8	Y	Y	4
<code>int (*A3)[3]</code>	Y	N	8	Y	Y	12			



Understanding Pointers & Arrays #2

Decl	<i>An</i>			<i>*An</i>			<i>**An</i>		
	Cmp	Bad	Size	Cmp	Bad	Size	Cmp	Bad	Size
<code>int A1[3]</code>	Y	N	12	Y	N	4	N	-	-
<code>int *A2[3]</code>	Y	N	24	Y	N	8	Y	Y	4
<code>int (*A3)[3]</code>	Y	N	8	Y	Y	12	Y	Y	4



Two-dimensional (Nested) Arrays

■ Declaration

$T \ A[R][C];$


- 2D array of data type T
- R rows, C columns

■ Array Size

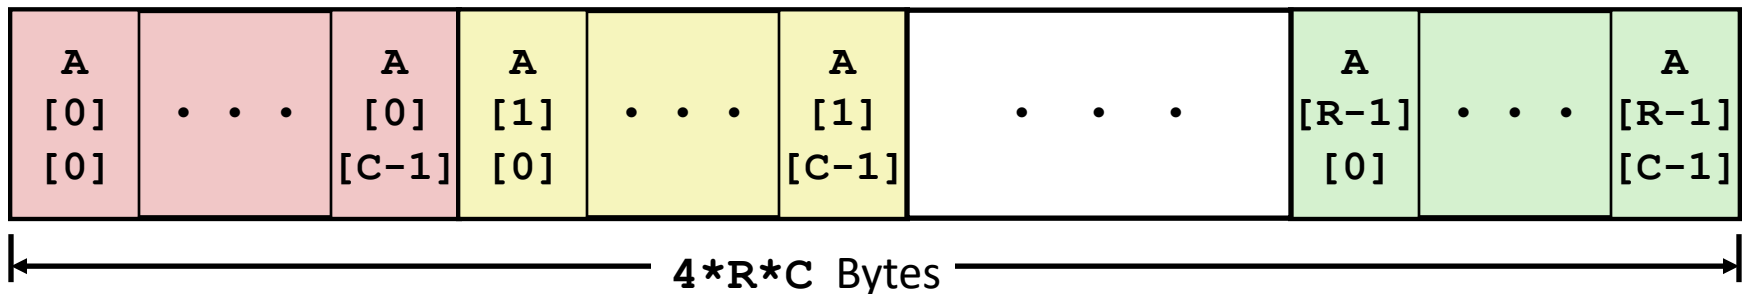
- $R * C * \text{sizeof}(T)$ bytes

■ Arrangement

- Row-Major Ordering




`int A[R][C];`



Array Example

```
#define XCOUNT 4
typedef int zip_dig[6];

zip_dig xian[XCOUNT] =
    {{7, 1, 0, 0, 4, 9 },
     {7, 1, 0, 0, 7, 1 },
     {7, 1, 0, 0, 7, 2 },
     {7, 1, 0, 0, 6, 9 }};
```




- “zip_dig xian[4]” equivalent to “int xian[4][6]”
 - Variable **xian**: array of 4 elements, allocated contiguously
 - Each element is an array of 6 **int**’s, allocated contiguously
- “Row-Major” ordering of all elements in memory

Nested Array Row Access

■ Row Vectors

- $A[i]$ is array of C elements
- Each element of type T requires $\text{sizeof}(T)$ bytes
- Starting address $A + i * (C * \text{sizeof}(T))$

```
int A[R][C];
```



Array Row Access Code

7	1	0	0	4	9	7	1	0	0	7	1	7	1	0	0	7	2	7	1	0	0	6	9
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

↑
xian

↑
xian[2]

```
int *get_xian_zip(int index)
{
    return xian[index];
}
```

```
# %rdi = index
leaq (%rdi,%rdi,2),%rax # 3 * index
leaq xian(,%rax,8),%rax # xian + (24 * index)
```

■ Row Vector

- **xian[index]** is array of 6 int's
- Starting address **xian+24*index**

■ Machine Code

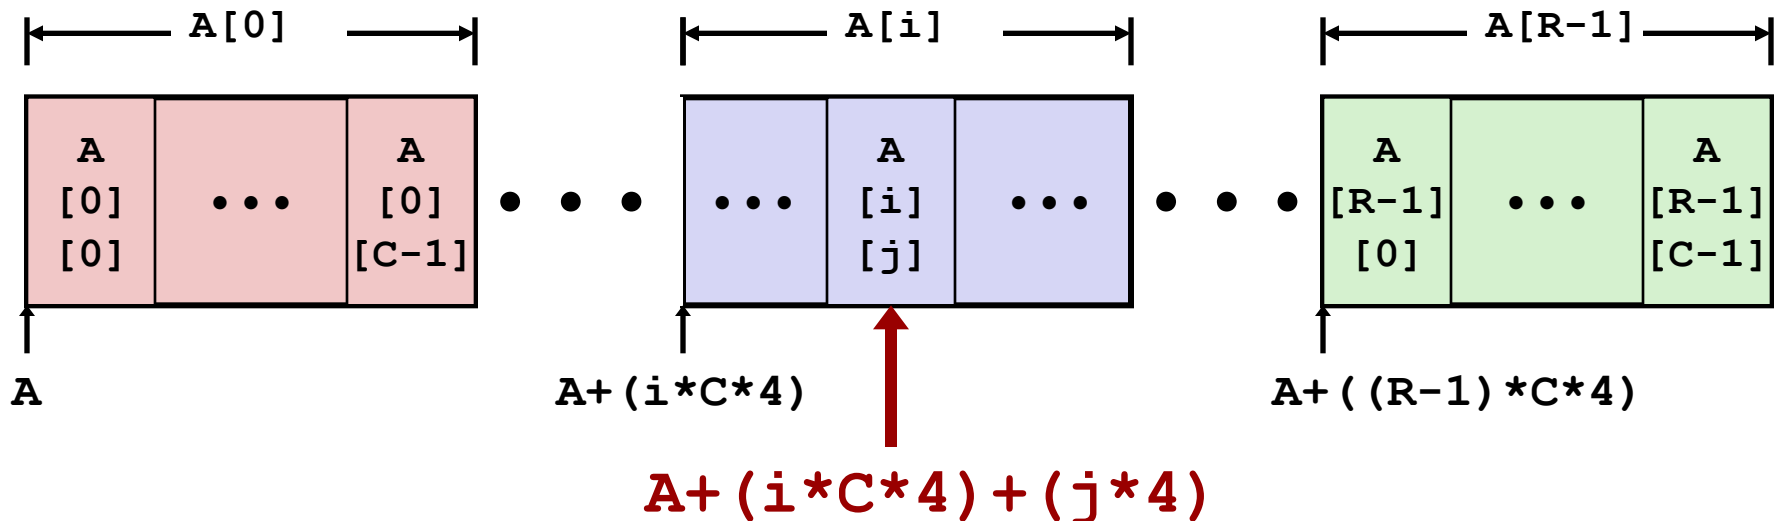
- Computes and returns address
- Compute as **xian + 8*(index+2*index)**

Nested Array Element Access


■ Array Elements

- $A[i][j]$ is element of type T , which requires $\text{sizeof}(T)$ bytes
- Address $A + i * (C * \text{sizeof}(T)) + j * \text{sizeof}(T)$
 $= A + (i * C + j) * \text{sizeof}(T)$

```
int A[R][C];
```



Array Element Access Code



```
int get_xian_digit(int index, int dig)
{
    return xian[index][dig];
}
```

```
leaq    (%rdi,%rdi,2), %rax    # 3*index
leaq    (%rsi,%rax,2), %rsi    # 6*index+dig
movl    xian(,%rsi,4), %eax    # xian + 4*(6*index+dig)
```

■ Array Elements


- `xian[index][dig]` is `int`
- Address: $\text{xian} + 24 \cdot \text{index} + 4 \cdot \text{dig}$
 $= \text{xian} + 4 \cdot (6 \cdot \text{index} + \text{dig})$

Multi-Level Array Example

```
zip_dig xjtu = { 7, 1, 0, 0, 4, 9 };
zip_dig thu  = { 1, 0, 0, 0, 8, 4 };
zip_dig pku  = { 1, 0, 0, 8, 7, 1 };
```


```
#define UCOUNT 3
int *univ[UCOUNT] = {thu, xjtu, pku};
```

- Variable `univ` denotes array of 3 elements
- Each element is a pointer
 - 8 bytes
- Each pointer points to array of `int`'s



Element Access in Multi-Level Array

```
int get_univ_digit
(size_t index, size_t digit)
{
    return univ[index][digit];
}
```



```
salq    $2, %rsi          # 4*digit
addq    univ(,%rdi,8), %rsi # p = univ[index] + 4*digit
movl    (%rsi), %eax       # return *p
ret
```

■ Computation

- Element access **Mem[Mem[univ+8*index]+4*digit]**
- Must do two memory reads
 - First get pointer to row array
 - Then access element within array



Array Element Accesses

Nested array

```
int get_xian_digit
(size_t index, size_t digit)
{
    return xian[index][digit];
}
```

Multi-level array

```
int get_univ_digit
(size_t index, size_t digit)
{
    return univ[index][digit];
}
```



Accesses looks similar in C, but address computations very different:

$\text{Mem}[\text{xian} + 24 * \text{index} + 4 * \text{digit}]$ $\text{Mem}[\text{Mem}[\text{univ} + 8 * \text{index}] + 4 * \text{digit}]$

N x N Matrix Code

■ Fixed dimensions

- Know value of N at compile time

```
#define N 16
typedef int fix_matrix[N][N];
/* Get element A[i][j] */
int fix_ele(fix_matrix A,
            size_t i, size_t j)
{
    return A[i][j];
}
```

■ Variable dimensions, explicit indexing

- Traditional way to implement dynamic arrays

```
#define IDX(n, i, j) ((i)*(n)+(j))
/* Get element A[i][j] */
int vec_ele(size_t n, int *A,
            size_t i, size_t j)
{
    return A[IDX(n,i,j)];
}
```

■ Variable dimensions, implicit indexing

- Added to language in 1999

```
/* Get element A[i][j] */
int var_ele(size_t n, int A[n][n],
            size_t i, size_t j) {
    return A[i][j];
}
```

16x16 Matrix Access

■ Array Elements

- `int A[16][16];`
- Address `A + i * (C * sizeof(int)) + j * sizeof(int)`
- `C = 16, sizeof(int) = 4`

```
/* Get element A[i][j] */
int fix_ele(fix_matrix A, size_t i, size_t j) {
    return A[i][j];
}
```

```
# A in %rdi, i in %rsi, j in %rdx
salq    $6, %rsi           # 64*i
addq    %rsi, %rdi         # A + 64*i
movl    (%rdi,%rdx,4), %eax # Mem[A + 64*i + 4*j]
ret
```

Matrix Access

■ Array Elements

- `size_t n;`
- `int A[n][n];`
- Address `A + i * (C * sizeof(int)) + j * sizeof(int)`
- `C = n, sizeof(int) = 4`
- Must perform integer multiplication

```
/* Get element A[i][j] */
int var_ele(size_t n, int A[n][n], size_t i, size_t j)
{
    return A[i][j];
}
```

```
# n in %rdi, A in %rsi, i in %rdx, j in %rcx
imulq    %rdx, %rdi          # n*i
leaq     (%rsi,%rdi,4), %rax  # A + 4*n*i
movl     (%rax,%rcx,4), %eax  # Mem[A + 4*n*i + 4*j]
ret
```

Example: Array Access

```
#include <stdio.h>

#define ZLEN 6
#define XCOUNT 4
typedef int zip_dig[ZLEN];

int main(int argc, char** argv) {
    zip_dig xian[XCOUNT] =
        {{7, 1, 0, 0, 4, 9},
         {7, 1, 0, 0, 7, 1 },
         {7, 1, 0, 0, 7, 2 },
         {7, 1, 0, 0, 6, 9 }};
    int *linear_zip = (int *) xian;
    int *zip2 = (int *) xian[2];
    int result =
        xian[0][0] +
        linear_zip[8] +
        *(linear_zip + 10) +
        zip2[1];
    printf("result: %d\n", result);
    return 0;
}
```

```
linux> ./array
```


Example: Array Access

```
#include <stdio.h>

#define ZLEN 6
#define XCOUNT 4
typedef int zip_dig[ZLEN];

int main(int argc, char** argv) {
    zip_dig xian[XCOUNT] =
        {{7, 1, 0, 0, 4, 9},
         {7, 1, 0, 0, 7, 1},
         {7, 1, 0, 0, 7, 2},
         {7, 1, 0, 0, 6, 9}};
    int *linear_zip = (int *) xian;
    int *zip2 = (int *) xian[2];
    int result =
        xian[0][0] +
        linear_zip[8] +
        *(linear_zip + 10) +
        zip2[1];
    printf("result: %d\n", result);
    return 0;
}
```

```
linux> ./array
result: 15
```



■ Arrays

- One-dimensional
- Multi-dimensional (nested)
- Multi-level


■ Structures

- Allocation
- Access
- Alignment

■ If we have time: Union

Structure Representation


```
struct rec {  
    int a[4];  
    size_t i;  
    struct rec *next;  
};
```



- **Structure represented as block of memory**
 - Big enough to hold all the fields
- **Fields ordered according to declaration**
 - Even if another ordering could be more compact
- **Compiler determines overall size + positions of fields**
 - In assembly, we see only offsets, not field names

Generating Pointer to Structure Member

```
struct rec {
    int a[4];
    size_t i;
    struct rec *next;
};
```



■ Generating Pointer to Array Element

- Offset of each structure member determined at compile time
- Compute as $r + 4 * idx$

```
int *get_ap
(struct rec *r, size_t idx)
{
    return &r->a[idx];
}
```


```
# r in %rdi, idx in %rsi
leaq (%rdi,%rsi,4), %rax
ret
```

Following Linked List #1

■ C Code

```
long length(struct rec*r) {
    long len = 0L;
    while (r) {
        len ++;
        r = r->next;
    }
    return len;
}
```

```
struct rec {
    int a[4];
    size_t i;
    struct rec *next;
};
```



Register	Value
%rdi	r
%rax	len

■ Loop assembly code


```
.L11:                # loop:
    addq    $1, %rax    # len ++
    movq    24(%rdi), %rdi    # r = Mem[r+24]
    testq   %rdi, %rdi    # Test r
    jne     .L11        # If != 0, goto loop
```

Following Linked List #2

■ C Code

```
void set_val
(struct rec *r, int val)
{
    while (r) {
        size_t i = r->i;
        // No bounds check
        r->a[i] = val;
        r = r->next;
    }
}
```

```
struct rec {
    int a[4];
    size_t i;
    struct rec *next;
};
```




Element i

Register	Value
%rdi	r
%rsi	val

```
.L11:                                # loop:
    movq    16(%rdi), %rax            # i = Mem[r+16]
    movl    %esi, (%rdi,%rax,4)      # Mem[r+4*i] = val
    movq    24(%rdi), %rdi          # r = Mem[r+24]
    testq   %rdi, %rdi              # Test r
    jne     .L11                    # if !=0 goto loop
```

Structures & Alignment

Unaligned Data




```
struct S1 {
    char c;
    int i[2];
    double v;
} *p;
```

Aligned Data

Primitive data type requires K bytes

Address must be multiple of K



Alignment Principles

■ Aligned Data

- Primitive data type requires B bytes
- Address must be multiple of B
- Required on some machines; advised on x86-64

■ Motivation for Aligning Data

- Memory accessed by (aligned) chunks of 4 or 8 bytes (system dependent)
 - Inefficient to load or store datum that spans quad word boundaries (8 bytes)
 - Inefficient to load or store datum that spans cache lines (64 bytes). Intel states should avoid crossing 16 byte boundaries.
 - Virtual memory trickier when datum spans 2 pages (4 KB pages)

■ Compiler

- Inserts gaps in structure to ensure correct alignment of fields

Specific Cases of Alignment (x86-64)

- **1 byte: char, ...**
 - no restrictions on address
- **2 bytes: short, ...**
 - lowest 1 bit of address must be 0_2
- **4 bytes: int, float, ...**
 - lowest 2 bits of address must be 00_2
- **8 bytes: double, long, char *, ...**
 - lowest 3 bits of address must be 000_2

Satisfying Alignment with Structures

Within structure:

Must satisfy each element's alignment requirement

Overall structure placement

Each structure has alignment requirement **K**

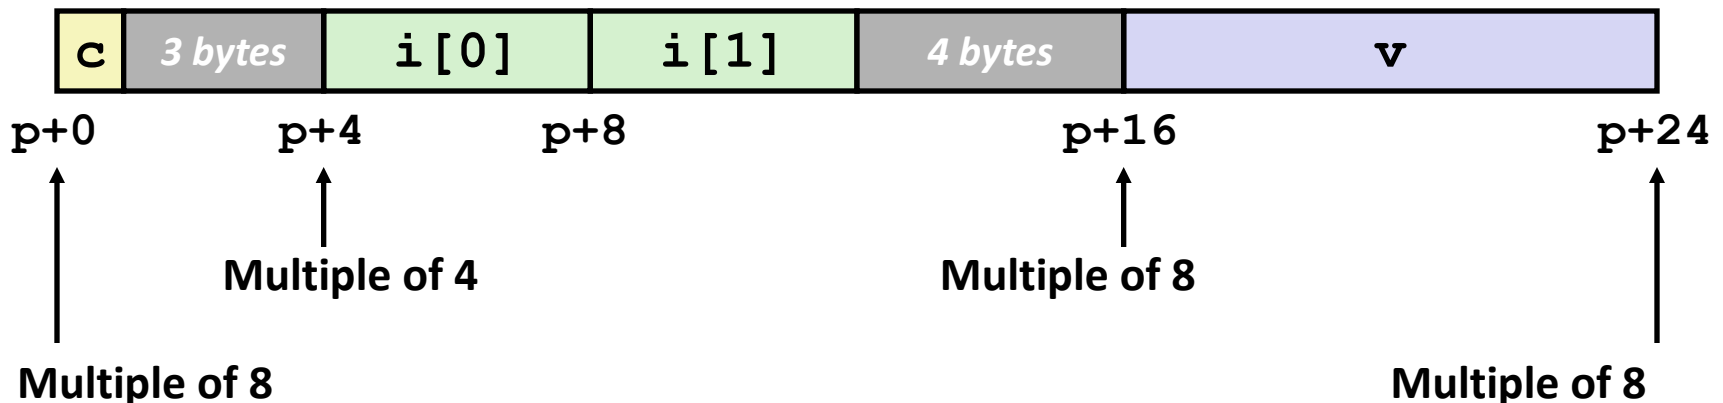
K = Largest alignment of any element

Initial address & structure length must be multiples of **K**

```
struct S1 {
    char c;
    int i[2];
    double v;
} *p;
```

Example:


K = 8, due to **double** element



Overall Alignment Requirement

For largest alignment requirement K
Overall structure must be multiple of K

```
struct S2 {  
    double v;  
    int i[2];  
    char c;  
} *p;
```




Multiple of K=8

Arrays of Structures

Overall structure length
multiple of K
Satisfy alignment requirement
for every element

```
struct S2 {  
    double v;  
    int i[2];  
    char c;  
} a[10];
```



Accessing Array Elements

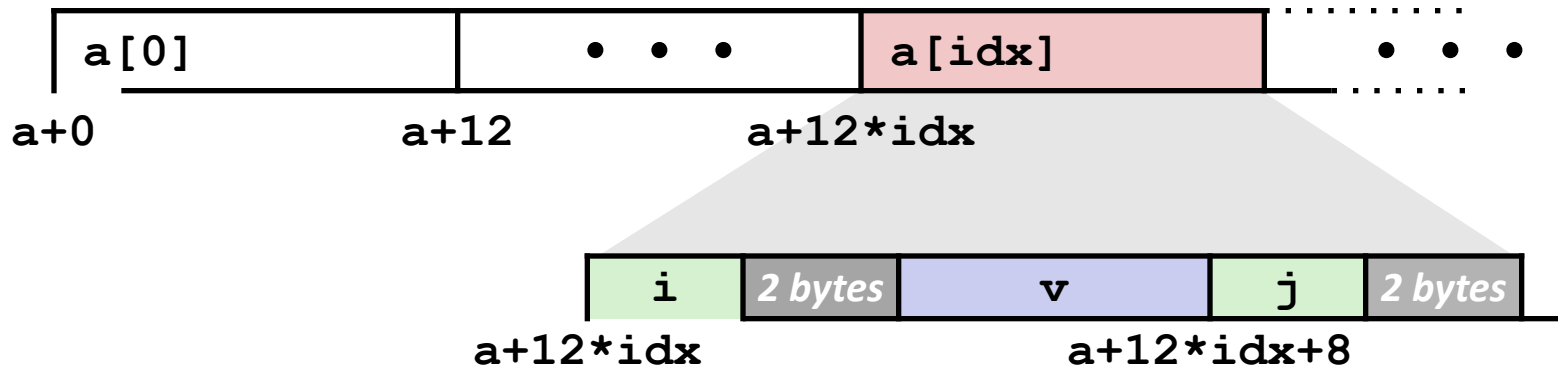
Compute array offset $12 * \text{idx}$

`sizeof(S3)`, including alignment spacers

Element `j` is at offset 8 within structure

Assembler gives offset `a+8`

Resolved during linking



```
struct S3 {
    short i;
    float v;
    short j;
} a[10];
```

```
short get_j(int idx)
{
    return a[idx].j;
}
```

```
# %rdi = idx
leaq (%rdi,%rdi,2),%rax # 3*idx
movzwl a+8(,%rax,4),%eax
```

Saving Space



Put large data types first


```
struct S4 {
    char c;
    int i;
    char d;
} *p;
```



```
struct S5 {
    int i;
    char c;
    char d;
} *p;
```

Effect (K=4)





■ Arrays

- One-dimensional
- Multi-dimensional (nested)
- Multi-level

■ Structures

- Allocation
- Access
- Alignment



■ If we have time: Union

Union Allocation

- Allocate according to largest element
- Can only use one field at a time


```
union U1 {
    char c;
    int i[2];
    double v;
} *up;
```

```
struct S1 {
    char c;
    int i[2];
    double v;
} *sp;
```



Using Union to Access Bit Patterns

```
typedef union {  
    float f;  
    unsigned u;  
} bit_float_t;
```



```
float bit2float(unsigned u)  
{  
    bit_float_t arg;  
    arg.u = u;  
    return arg.f;  
}
```

Same as (float) u?



```
unsigned float2bit(float f)  
{  
    bit_float_t arg;  
    arg.f = f;  
    return arg.u;  
}
```

Same as (unsigned) f?

Using Unions as Sum Types

```
typedef union {
    float f;
    unsigned u;
} num_t;
```

```
typedef struct {
    bool is_float;
    num_t val;
} value_t;
```



(technically `is_float` only takes 1 byte and then there's 3 bytes of padding)

Summary

■ Arrays

- Elements packed into contiguous region of memory
- Aligned to satisfy every element's alignment requirement
- Pointer to first element
- Use index arithmetic to locate individual elements
- No bounds checking

■ Structures

- Elements packed into single region of memory
- Possible require internal and external padding to ensure alignment
- Access using offsets determined by compiler

■ Unions

- Overlay declarations
- Way to circumvent type system